



YOUR THESIS TITLE HERE

Your Full Name

Supervisor: Your Supervisor Name

School of Mathematics and Statistics
UNSW Sydney

July 2026

SUBMITTED IN PARTIAL FULFILMENT OF THE REQUIREMENTS OF THE DEGREE OF
MASTER OF STATISTICS

Plagiarism statement

I declare that this thesis is my own work, except where acknowledged, and has not been submitted for academic credit elsewhere.

I acknowledge that the assessor of this thesis may, for the purpose of assessing it:

- Reproduce it and provide a copy to another member of the University; and/or,
- Communicate a copy of it to a plagiarism checking service (which may then retain a copy of it on its database for the purpose of future plagiarism checking).

I certify that I have read and understood the University Rules in respect of Student Academic Misconduct, and am aware of any potential plagiarism penalties which may apply.

By signing this declaration I am agreeing to the statements and conditions above.

Signed: _____

Date: _____

Abstract

Write your abstract here.

Contents

Chapter 1	Literature Review	1
1.1	Foundations of Belief and Updating & Variational Inference	1
1.2	The Shift to Predictively Oriented Posteriors	7
1.3	Particle-Based Inference, MMD, and Misspecification	10
1.3.1	Stein varational gradient descent (SVGD)	11
1.3.2	Deep Ensembles: Omitting the Dispersion Force	12
1.3.3	MMD of PrO	15
Chapter 2	Other notes for later chapters:	19
References		22

CHAPTER 1

Literature Review

1.1 Foundations of Belief and Updating & Variational Inference

At the core of statistical inference is the challenge of mathematically capturing the structure and uncertainty inherent in observed data.

Suppose that we have a statistical model $P_\theta \in \mathcal{P}$, where $\mathcal{P} = \{P_\theta : \theta \in \Theta\}$, with the parameters $\theta \in \Theta$ for the data points $x_{1:n} = \{x_1, x_2, \dots, x_n\}$ with a theoretical and often unknown data generating process (DGP) $D(x)$. The parameters θ index the candidate distributions, with their specific values typically encoding information about the data itself. In frequentist statistics these parameters θ are treated as fixed but unknown constants, however, when dealing with the Bayesian approach, the parameters are treated as random variables / distributions. In the most standard Bayes we have the posterior density

$$\pi_B(\theta | x_{1:n}) = \frac{\pi(\theta)L(x_{1:n} | \theta)}{m(x_{1:n})}.$$

Where $\pi(\theta)$ is our prior knowledge of the parameters θ , $L(x_{1:n} | \theta)$ is the likelihood function representing the probability of the observed data under those parameters, and $\pi_B(\theta | x_{1:n})$ is the pdf of the distribution $\theta | x_{1:n}$. The B here indicating that this posterior is the standard Bayesian posterior. Here $m(x_{1:n}) = \int_{\Theta} L(x_{1:n} | \theta) d\pi(\theta)$ is the marginal normalising constant of the distribution, whose computation often makes the distribution intractable.

Bayesian inference is conceptually and practically appealing as, unlike frequentist statistical methods, we are able to allow our previous knowledge, specific domain expertise and experts in the subject matter to further inform our distribution. This injection of previous knowledge comes in the form of a carefully specified prior $\pi(\theta)$. Naturally, this comes with the potential risk of poor prior choice.

A further advantage is that Bayesian posteriors produce belief distributions (as opposed to point estimates, such as the direct outputs of MLE) over the parameters of interest $\theta \in \Theta$.

This carries two related benefits. First, prediction propagates parameter uncertainty rather than conditioning on a single value of θ ; where the conventional Bayesian predictive density

for an unobserved x_{new} integrates the model over the posterior,

$$p(x_{\text{new}} | x_{1:n}) = \int_{\Theta} p(x_{\text{new}} | \theta) \pi_B(\theta | x_{1:n}) d\theta. \quad (1.1)$$

Second, any quantity of interest is recovered as a functional of the posterior; for instance, the posterior mean $\hat{\theta}$ is recovered as $\int_{\Theta} \theta \pi_B(\theta | x_{1:n}) d\theta$. Because inference averages over θ instead of fixing it at a single *best* value (*best* referring to the best available point estimate of θ , e.g., the MLE or MAP), the resulting predictions are less prone to the over-confidence of a plug-in point estimate.

Moving to a more generalised Bayes, we can recast inference as an optimisation-centric approach [Knoblauch et al. \(2022\)](#). Generalised Bayes replaces the likelihood with a loss function raised to a negative an exponential.

$$\begin{aligned} \pi_w(\theta | x_{1:n}) &= \frac{\exp\left\{-w \sum_{i=1}^n \ell(\theta, x_i)\right\} \pi(\theta)}{\int_{\Theta} \exp\left\{-w \sum_{i=1}^n \ell(\theta, x_i)\right\} \pi(\theta) d\theta} \\ &\propto \exp\left\{-w \sum_{i=1}^n \ell(\theta, x_i)\right\} \pi(\theta) \end{aligned} \quad (1.2)$$

where, as before, the denominator is the normalising constant; without it, we characterise the distribution only up to proportionality. There also exists a learning rate or loss scale $w \in \mathbb{R}_{>0}$. This posterior is the Gibbs posterior introduced by [Bissiri et al. \(2016\)](#).

It is important to note here that from this posterior we can recover our previous standard Bayes posterior π_B in the specific case with $\ell(\theta, x_i) = -\log p(x_i | \theta)$ and the learning rate $w = 1$. The loss $-\log p(x_i | \theta)$ is the self-information (surprisal) of each observation, so this choice penalises the model exactly by how 'surprised' it is by the data. Following this we get

$$\begin{aligned} \pi_w(\theta | x_{1:n}) &\propto \exp\left\{-w \sum_{i=1}^n \ell(\theta, x_i)\right\} \pi(\theta) \\ &= \exp\left\{\sum_{i=1}^n \log p(x_i | \theta)\right\} \pi(\theta) \\ &= \left(\prod_{i=1}^n p(x_i | \theta)\right) \pi(\theta) \\ &= L(x_{1:n} | \theta) \pi(\theta) \end{aligned}$$

Furthermore, the normalising constant recovers the marginal likelihood in the same manner. Hence we recover $\pi_B(\theta | x_{1:n})$. So standard Bayes is not a separate paradigm but a single specific instance of a more generalised approach where the loss is the log-loss and $w = 1$ [Knoblauch et al. \(2022\)](#).

The learning rate w

Looking at the learning rate w here, we can see this as a weighting between the prior and the data, and a guard against model misspecification (assuming the prior is well specified). It acts as an effective-sample-size multiplier, effectively updating the prior with only wn data points. Observing the limiting behaviour we see as

$$w \rightarrow 0^+ : \pi_w \rightarrow \pi \text{ (prior)}; \quad w \rightarrow \infty : \pi_w \rightarrow \delta_{\hat{\theta}}, \quad \hat{\theta} = \arg \min_{\theta} \sum_i \ell(\theta, x_i).$$

The effect of this learning rate is exact and analytic and *does* affect the final distribution of the posterior. Here w is part of the target distribution itself, not of any algorithm used to approximate it. This should be held separate from the step size of an optimiser, which (for a non-analytic solver) governs the trajectory toward a fixed target rather than the target. This distinction is flagged now because this property is altered later. Once the posterior is itself defined as the solution of a particle optimisation (PVI), the data-weight and the optimiser's effect on the realised distribution no longer cleanly separate. This is further discussed in Section ??.

Coherence and the loss-based view

[Bissiri et al. \(2016\)](#) do not posit the exponentiated-loss form; they derive it from a *coherence* requirement on how beliefs may be updated. Coherence demands that the update be self-consistent under partitioning of the data: updating on a first batch and then on a second must yield the same belief as updating on the pooled sample at once, so that the order and grouping of the data cannot change our conclusions. An update with (x_1, x_2) together or $\{(x_1), (x_2)\}$ separately yields the same result. This coherence property can be formally expressed as

$$\psi[l(\theta, x_2), \psi\{l(\theta, x_1), \pi(\theta)\}] = \psi\{l(\theta, x_2) + l(\theta, x_1), \pi(\theta)\},$$

where the function ψ is the update function such that in the standard case

$$\pi(\theta | x) \equiv \psi\{l(\theta; x), \pi(\theta)\}.$$

For an additive loss they show the unique update satisfying this is the Gibbs posterior above. This matters going forward because coherence is a property of the loss: the self-information (log) loss satisfies it, but as we adopt richer, predictive and discrepancy-based losses later, coherence is no longer automatic and must be checked rather than assumed.

One may often only be interested in a simple, low-dimensional statistic, such as a population mean or median, which provides further motivation for a loss-based posterior. In standard Bayesian frameworks, estimating these simple statistics typically requires modelling the entire data generating process (as in π_B). However, issues arise when the true form of the data distribution is unknown. As highlighted by [Bissiri et al. \(2016\)](#), if there is no true parameter θ

that perfectly connects to the data via a known density, a scenario referred to as the $\mathcal{M}$ -open case [Bernardo and Smith \(1994\)](#), the standard Bayesian update struggles. Furthermore, some parameters of interest simply do not directly index any standard parametric family of density functions.

Instead of relying on a full probability model, a loss function can connect a parameter to the data directly. For example, suppose we have θ as our parameter of interest representing the population median. We can update our beliefs about θ directly by employing the absolute deviation loss, $\ell(\theta, x_i) = |\theta - x_i|$, entirely bypassing the need to assume the data itself is Laplace distributed [Bissiri et al. \(2016\)](#). However, it is important to note that absolute-error-loss inference can still succumb to misspecification; a robust loss limits, but does not entirely remove, sensitivity to the true data generating process (DGP).

In our setting later on, we nonetheless retain the full posterior rather than focusing on a single functional, as doing so is essential for capturing parameter uncertainty and propagating it faithfully into our predictive distributions.

The optimisation-centric view and the Rule of Three

The generalised Bayes posterior π_w has an equivalent expression focusing more on the optimisation perspective. [Knoblauch et al. \(2022\)](#) show that the exact posterior can equally be obtained as the solution of an infinite-dimensional optimisation problem over the space $\mathcal{P}(\Theta)$ of *all* probability measures on Θ ; a representation they trace to Csiszár (1975), Donsker & Varadhan (1975), and Zellner (1988), where

$$\pi_B(\theta \mid x_{1:n}) = \arg \min_{q \in \mathcal{P}(\Theta)} \left\{ \mathbb{E}_q \left[\sum_{i=1}^n \ell(\theta, x_i) \right] + \text{KL}(q \parallel \pi) \right\}, \quad \ell(\theta, x_i) = -\log p(x_i \mid \theta).$$

Replacing the log-loss with a general ℓ (and absorbing the loss scale w into ℓ) returns exactly the generalised Gibbs posterior of the previous section; [Bissiri et al. \(2016\)](#) show the minimiser exists and coincides with π_w ¹ whenever the normaliser $Z = \int_{\Theta} \exp\{-\sum_i \ell(\theta, x_i)\} \pi(\theta) d\theta$ is finite. The objective can be interpreted as regularised empirical risk minimisation; it functions in two parts, minimising the expected loss over the data while the KL term pulls the solution toward the prior. These are both important points that will be further discussed later on. It should be noted that here it is the KL term which causes the optimisation problem to converge to a distribution, without it the problem collapses to a point mass $\delta_{\hat{\theta}}$ at the empirical risk minimiser $\hat{\theta} = \arg \min_{\theta} \sum_i \ell(\theta, x_i)$. This can be seen as we derive our previous bayes posteriors from this optimisation results.

¹[Knoblauch et al. \(2022\)](#) do not notationally distinguish the standard and generalised Bayesian posteriors, writing q_B^* for both; here we keep them separate as π_B and π_w .

Letting $\ell_n(\theta) = \sum_{i=1}^n \ell(\theta, x_i)$ denote the cumulative loss, the normalising constant Z and the normalised Gibbs posterior $\pi_w(\theta)$ are defined as:

$$Z = \int_{\Theta} \pi(\theta) e^{-\ell_n(\theta)} d\theta \quad \text{and} \quad \pi_w(\theta) = \frac{\pi(\theta) e^{-\ell_n(\theta)}}{Z}.$$

The objective then collapses into a single Kullback-Leibler (KL) divergence, as

$$\begin{aligned} \mathbb{E}_q \left[\sum_{i=1}^n \ell(\theta, x_i) \right] + \text{KL}(q \parallel \pi) &= \int_{\Theta} q(\theta) \ell_n(\theta) d\theta + \int_{\Theta} q(\theta) \log \frac{q(\theta)}{\pi(\theta)} d\theta \\ &= \int_{\Theta} q(\theta) \left[\log \frac{q(\theta)}{\pi(\theta)} + \ell_n(\theta) \right] d\theta \\ &= \int_{\Theta} q(\theta) \log \frac{q(\theta)}{\pi(\theta) e^{-\ell_n(\theta)}} d\theta \\ &= \int_{\Theta} q(\theta) \log \frac{q(\theta)}{Z \pi_w(\theta)} d\theta \\ &= \text{KL}(q \parallel \pi_w) - \log Z. \end{aligned}$$

Since $\log Z$ is independent of q and $\text{KL}(q \parallel \pi_w) \geq 0$ with equality if and only if $q = \pi_w$, the minimiser over $\mathcal{P}(\Theta)$ is exactly the Gibbs posterior

$$\arg \min_{q \in \mathcal{P}(\Theta)} \left\{ \mathbb{E}_q \left[\sum_{i=1}^n \ell(\theta, x_i) \right] + \text{KL}(q \parallel \pi) \right\} = \pi_w(\theta) = \frac{\pi(\theta) \exp \{ -\sum_{i=1}^n \ell(\theta, x_i) \}}{Z}.$$

Specialising to the log-loss $\ell(\theta, x_i) = -\log p(x_i \mid \theta)$ yields $e^{-\ell_n(\theta)} = \prod_{i=1}^n p(x_i \mid \theta) = L(x_{1:n} \mid \theta)$ alongside the marginal likelihood $Z = \int_{\Theta} \pi(\theta) L(x_{1:n} \mid \theta) d\theta = m(x_{1:n})$. Substituting these components gives us

$$\pi_w(\theta) = \frac{\pi(\theta) L(x_{1:n} \mid \theta)}{m(x_{1:n})} = \pi_B(\theta \mid x_{1:n}),$$

which recovers the standard Bayesian posterior. For more information and a formal proof regarding the collapse to a point mass, see [Knoblauch et al. \(2022\)](#).

Furthermore, we can see now that by using the standard Bayesian posterior, we are equivalently committing to this specific optimisation problem. As [Knoblauch et al. \(2022\)](#) point out, similar to generalised Bayes, using this standard posterior is only the right tool if minimising this objective actually reflects the overarching goal of the inference task. If the standard log-loss does not align with those practical goals, this framework justifies replacing it with a custom loss function. This transformation to an optimisation problem is what we will focus on throughout this thesis, and from here on out.

Taking inference as an optimisation, the standard posterior is seen to fix three choices that need not be fixed: a loss ℓ , the divergence KL measuring departure from the prior, and the feasible set $\mathcal{P}(\Theta)$ optimised over. [Knoblauch et al. \(2022\)](#) free all three, defining the *Rule of*

Three (RoT) posterior:

$$\pi_{(\ell, D, \Pi)}^* := \arg \min_{q \in \Pi} \left\{ \mathbb{E}_q \left[\sum_{i=1}^n \ell(\theta, x_i) \right] + D(q \| \pi) \right\}, \quad (1.3)$$

specified by a loss ℓ , a divergence D , and a feasible set $\Pi \subseteq \mathcal{P}(\Theta)^2$. Standard Bayes is the choice $\pi_{(-\log p, \text{KL}, \mathcal{P}(\Theta))}^*$ and generalised Bayes is $\pi_{(\ell, \text{KL}, \mathcal{P}(\Theta))}^*$; the same template also recovers tempered and PAC-Bayesian posteriors. The three arguments answer three assumptions that justify standard Bayes, though these may be problematic in practice: a well-specified prior (P), a well-specified likelihood (L), and unlimited computation (C). The divergence D controls how, and how strongly, the prior is enforced (P); the loss ℓ controls robustness to model misspecification (L); and the feasible set Π controls tractability (C).

An issue arises in this setup where, computationally, this optimisation becomes quite intractable over an infinite-dimensional space; the third argument in our objective, Π , defined as the space of all possible probability measures $\mathcal{P}(\Theta)$, is infinite-dimensional. Variational Inference (VI) helps address this by restricting this third argument to a finite-dimensional family $\mathcal{Q} = \{q(\theta | \kappa) : \kappa \in K\}$. VI shrinks an impossible infinite-dimensional search down to a tractable finite one. Thus, while VI is an approximation of the true posterior, it is still shown to be grounded by the same optimisation problem but constrained, yielding

$$\pi_{\text{VI}}^* = \pi_{(\ell, \text{KL}, \mathcal{Q})}^*.$$

Standard VI is then not a separate approximation principle but the *best $\mathcal{Q}$ -constrained solution of the very same objective* (Knoblauch et al., 2022, Thm. 2). The tractable case in which $\Pi = \mathcal{Q}$ is a variational family and D may differ from KL is what they call *Generalised Variational Inference* (GVI). This shifts the role of $\mathcal{Q}$; rather than treating inference in $\mathcal{Q}$ as an approximation to π_B , we may take the objective itself as the design target and construct belief distributions in $\mathcal{Q}$ with the properties we want from the outset.

Variational and stochastic variational inference

With both the standard and generalised framework, typically the resulting posterior is intractable due to the normalising constant $p(x_{1:n})$ (There are cases where it is tractable in simple examples or Conjugate Priors, but simplifying a model to get this result often leads to larger model misspecification). This evidence represents the total probability of observing the data marginalised over all possible parameter configurations, calculated as $Z = p(x_{1:n}) = \int_{\Theta} p(x_{1:n} | \theta) \pi(\theta) d\theta$.

Variational inference attempts to circumvent this by restricting the search to a tractable family of distributions $\mathcal{Q}$ and selecting the distribution that is closest to the true posterior in

²Notation: Knoblauch et al. (2022) et al. write the RoT posterior as $P(\ell, D, \Pi)$. Here it is $\pi_{(\ell, D, \Pi)}^*$

Kullback-Leibler divergence. The ideal variational distribution is therefore

$$q^* = \arg \min_{q \in \mathcal{Q}} \text{KL}(q(\theta) \| \pi_B(\theta | x_{1:n})).$$

However, by expanding the KL divergence we see that the divergence cannot be minimised directly because its evaluation depends on the intractable evidence we are trying to avoid, yielding

$$\begin{aligned} \text{KL}(q(\theta) \| \pi_B(\theta | x_{1:n})) &= \int_{\Theta} q(\theta) \log \frac{q(\theta)}{\pi_B(\theta | x_{1:n})} d\theta \\ &= \int_{\Theta} q(\theta) \log \frac{q(\theta) p(x_{1:n})}{p(x_{1:n} | \theta) \pi(\theta)} d\theta \\ &= \int_{\Theta} q(\theta) [\log q(\theta) - \log \pi(\theta) - \log p(x_{1:n} | \theta)] d\theta + \log p(x_{1:n}) \\ &= \text{KL}(q(\theta) \| \pi(\theta)) - \mathbb{E}_q[\log p(x_{1:n} | \theta)] + \log p(x_{1:n}). \end{aligned}$$

By rearranging these terms, we can express the log evidence as the sum of two distinct components,

$$\log p(x_{1:n}) = (\mathbb{E}_q[\log p(x_{1:n} | \theta)] - \text{KL}(q(\theta) \| \pi(\theta))) + \text{KL}(q(\theta) \| \pi_B(\theta | x_{1:n})).$$

The bracketed term is known as the Evidence Lower Bound (ELBO). Because the KL divergence is strictly non-negative, the ELBO strictly lower-bounds the true log evidence. Furthermore, since the log evidence $p(x_{1:n})$ is a constant with respect to the variational distribution q , minimising the KL divergence to the exact posterior is mathematically identical to maximising this lower bound, formally defined as

$$\text{ELBO}(q) = \mathbb{E}_q[\log p(x_{1:n} | \theta)] - \text{KL}(q(\theta) \| \pi(\theta)).$$

Classical coordinate-ascent variational inference optimises this ELBO with a full sweep over the dataset at each iteration, which quickly becomes computationally prohibitive for large datasets noted in [Blei et al. \(2017\)](#); [Hoffman et al. \(2013\)](#). To address this, Stochastic Variational Inference [Hoffman et al. \(2013\)](#) applies stochastic optimisation by forming noisy, unbiased natural-gradient estimates of the ELBO from random data minibatches. The defining limitation is that they fix a parametric family $\mathcal{Q}$ in advance. The quality of the final approximation can never exceed the expressiveness of this chosen family.

1.2 The Shift to Predictively Oriented Posteriors

The Rule-of-three optimisation-centric view outlined earlier allows for the free deviation of the loss ℓ , the divergence D , and the feasible set $\Pi \subseteq \mathcal{P}(\Theta)$; however, leaving the ‘parameter-centric focus’ untouched. In every posterior above we ultimately learn a belief distribution over the parameter θ , and prediction is a downstream by-product recovered through (1.1). While

Aitchison (1975) proves that the Bayesian predictive (1.1) outperforms a single fixed $\hat{\theta}$, and explicitly argues that predictive accuracy is the ultimate test of a model, the traditional Bayesian update remains structurally blind to this predictive goal during the learning phase. Predictively Oriented (PrO) inference (McLatchie et al., 2025) makes the predictive distribution itself the object of inference. It expresses parameter uncertainty as a consequence of predictive ability, rather than learning θ and treating prediction merely as a byproduct. This framework seeks to obtain “the most desirable aspects of both parameter inference and density estimation” (McLatchie et al., 2025). It sits alongside the closely related predictive variational inference of Lai and Yao (2024) and the MMD construction of Shen et al. (2024), which is taken up in Section ??.

Looking deeper at the parameter-centric posteriors of Section 1 and their result for prediction, when the model is well specified the DGP $D(x)$ is recoverable by some $P_{\theta_*} \in \mathcal{P}(\mathcal{M})$ -closed as discussed earlier). We see that in posterior concentration, where asymptotically, as n grows, q contracts onto θ_* and the predictive P_q becomes well calibrated. Under misspecification, where no single θ reproduces the data ($\mathcal{M}$ -open), concentration becomes actively harmful. McLatchie et al. (2025) show that both classical and generalised Bayes posteriors nonetheless concentrate as $n \rightarrow \infty$ (Miller, 2021), so that the posterior predictive

$$P_q = \int_{\Theta} P_{\theta} dq(\theta)$$

ceases to behave as a genuine combination of distinct models and degenerates to the prediction of a single pseudo-true parameter. The model-averaging content of (1.1) is lost precisely when it is most needed. The PrO posterior changes this target.

Motivating example image to be made or found to be added

Scoring rules as the loss

To orient inference towards prediction the generic loss ℓ of Section 1 is instantiated as a *strictly proper scoring rule* S . A scoring rule $S(P, x)$ returns a penalty for having issued the predictive distribution P when the outcome x is subsequently observed. It is proper if reporting the true distribution is optimal in expectation,

$$\mathbb{E}_{x \sim P_*}[S(P_*, x)] \leq \mathbb{E}_{x \sim P_*}[S(P, x)] \quad \forall P \in \mathcal{P}(\mathcal{Y}),$$

and strictly proper if equality holds only at $P = P_*$. The π_B score is $S(P, x) = -\log p(x)$, which is strictly proper and recovers the self-information loss of Section 1; this is further discussed in Gneiting and Raftery (2007).

Scoring the ensemble, not the particle

The distinction is *where* the scoring rule is evaluated relative to the averaging over q . The generalised posterior of Section 1, written here with S in place of ℓ ,

$$\pi_w = \arg \min_{q \in \mathcal{P}(\Theta)} \left\{ w \sum_{i=1}^n \int_{\Theta} S(P_{\theta}, x_i) q(\theta) d\theta + \text{KL}(q \parallel \pi) \right\},$$

evaluates the score *inside* the integral: it scores each P_{θ} in isolation and averages, so it rewards a q concentrated on parameters that *individually* predict well and penalises any single θ that is confidently wrong. The PrO posterior instead scores the marginal predictive P_q itself,

$$\pi_{\text{PrO}} = \arg \min_{q \in \mathcal{P}(\Theta)} \left\{ w \sum_{i=1}^n S(P_q, x_i) + \text{KL}(q \parallel \pi) \right\}, \quad P_q := \int_{\Theta} P_{\theta} q(\theta) d\theta,$$

³moving the focus from individual parameters to the collective ensemble.

The two objectives are ordered by Jensen's inequality. For the logarithmic score the map $P \mapsto S(P, x) = -\log p(x)$ is convex in the mixture density $p_q(x) = \int_{\Theta} p_{\theta}(x) dq(\theta)$, which is linear in q , so

$$S(P_q, x_i) = -\log \int_{\Theta} p_{\theta}(x_i) dq(\theta) \leq - \int_{\Theta} \log p_{\theta}(x_i) dq(\theta) = \int_{\Theta} S(P_{\theta}, x_i) dq(\theta),$$

with equality only if every ensemble member assigns x_i the same density. The mixture predicts at least as well as its average member, and the slack is exactly the benefit of model averaging. The generalised posterior minimises the upper bound and so has no incentive to spread mass across complementary models; the PrO posterior minimises the quantity actually of interest and is therefore free to retain a diverse ensemble whenever diversity sharpens the mixture. This is the ensemble advantage that posterior collapse destroys.

The stationarity condition

Because P_q depends on q , the PrO objective is no longer linear in q and admits no closed-form Gibbs solution. Taking the first variation under the constraint $\int dq = 1$, and writing $\frac{\delta S(P_q, x_i)}{\delta q}(\theta)$ for the functional derivative of the score at P_q projected onto P_{θ} , stationarity requires

$$w \sum_{i=1}^n \frac{\delta S(P_q, x_i)}{\delta q}(\theta) + \log \frac{q(\theta)}{\pi(\theta)} + 1 = \text{const},$$

which rearranges to the implicit fixed point

$$\pi_{\text{PrO}}(\theta) \propto \pi(\theta) \exp \left\{ -w \sum_{i=1}^n \frac{\delta S(P_q, x_i)}{\delta q} \Big|_{q=\pi_{\text{PrO}}}(\theta) \right\}.$$

³In the paper they use $\frac{\lambda_n}{n}$, which serves the exact same function as the previously defined w .

The defining feature is self-consistency: π_{PrO} appears on both sides, because the effective per-sample loss $\frac{\delta S(P_q, x_i)}{\delta q}$ is itself a functional of the posterior being solved for. Contrast Section 1, where the Gibbs loss $S(P_\theta, x_i)$ is independent of q and the exponential-family minimiser is explicit. The PrO posterior has no such closed form and must be reached by iteration — which is what motivates the particle and gradient-flow machinery of Section ??; McLatchie et al. (2025) obtain it by evolving particles under mean-field Langevin dynamics.

Predictive dominance

The reward for changing the target is a predictive guarantee the parameter-centric posteriors cannot match. McLatchie et al. (2025) establish that the PrO posterior *predictively dominates* both classical and generalised Bayes posterior predictives, converging — up to logarithmic factors — to the predictively optimal model average irrespective of specification. When the model is well specified this coincides with ordinary posterior concentration, so nothing is lost; when it is misspecified the PrO posterior does not collapse but stabilises at a predictively optimal limit, whose residual, irreducible uncertainty admits a direct reading as the degree of model misspecification. Uncertainty becomes a consequence of predictive ability rather than an artefact of the parameter count.

Gibbs minimises the average of scores; PrO minimises the score of the average.

1.3 Particle-Based Inference, MMD, and Misspecification

At the end of Section 1 we saw a structural limitation with standard Variational and stochastic variational inference. That is, to make our approximation tractable we fix a parametric family $\mathcal{Q}$ in advance. However, by setting an approximating distribution we limit how rich our expression can be to the maximum the selected family permits; under model misspecification, where the best description of the data may be multimodal or heavy-tailed, this restriction is what degrades the approximation.

With particle-based methods (or particle approximations), we remove this constraint. Instead of parameterising q (where q is our variational approximation of the true target distribution π_w as noted in section 1 (elbo)) and optimising q 's parameters, we use a cloud of M particles $\{\theta_1, \dots, \theta_M\}$ with the empirical measure $q \approx \frac{1}{M} \sum_j \delta_{\theta_j}$ and transport the particles directly. By transitioning to particle-based methods, we redefine the feasible set Π within the Rule of Three framework, replacing the rigid parametric $\mathcal{Q}$ with the infinitely flexible space of empirical measures.

At the core of particle-based methods is an interacting physical system. To attempt to achieve the target, we apply forces that iteratively nudge our particles through the parameter space.

We can conceptualise the movement and final distribution of these particles as being governed by three primary forces:

1. The pull of the data (The Likelihood/Loss): A driving force that drags particles toward parameter values that minimise the empirical risk or self-information loss.
2. The anchor of the prior: A regularising force that pulls particles back toward our initial belief distribution $\pi(\theta)$, preventing overfitting.
3. The dispersion force (Repulsion/Noise): A mechanism that scatters the particles. This prevents the ensemble from collapsing into a single point estimate (MAP) and ensures the particles spread out to accurately map the uncertainty and variance of the target distribution.⁴

Particle-based inference methods generally manage the trade-off between mode-seeking behaviour and particle diversity through three primary mechanisms.

We will focus on two complementary lineages; the deterministic Stein flow of Liu and Wang (2016), and the Wasserstein-gradient-flow view of Wild et al. (2023), which unifies deep ensembles, Langevin dynamics and variational inference under a single functional. This section reviews both, and reconstructs the second on the toy problems of Wild et al. to expose the machinery our own method inherits.

1.3.1 Stein variational gradient descent (SVGD)

Stein variational gradient descent (Liu and Wang, 2016) approximates a target $p(\theta)$ by deterministically transporting particles along the steepest descent of $\text{KL}(q||p)$ within the unit ball of a Reproducing-kernel hilbert space. **Explain the difference between L2 and RKHS and the norm perhaps**




Figure 1.1: An Empirical Kernel Mean Embedding mapping discrete particles into a Reproducing Kernel Hilbert Space (RKHS). The bandwidth (σ) visually dictates the strictness of the RKHS smoothness penalty, illustrating how isolated data points are mathematically transformed into the continuous, interacting field that drives SVGD’s repulsive force.

The three conceptual forces are directly seen in the SVGD update equation. Let our ensemble of M particles at iteration t be $\{\theta_j^{(t)}\}_{j=1}^M$. The particles are updated deterministically via

⁴Forces 2 and 3 are not completely independent as we will see in Deep Ensembles

a step size ε giving an update step

$$\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} + \varepsilon \phi^*(\theta_j^{(t)}). \quad (1.4)$$

The optimal perturbation direction $\phi^*(\theta_j)$ is restricted to a Reproducing Kernel Hilbert Space (RKHS) governed by a positive-definite kernel $\kappa(\cdot, \cdot)$. It is computed as the empirical average over the entire particle ensemble:

$$\phi^*(\theta_j) = \frac{1}{M} \sum_{k=1}^M \left[\underbrace{-w \sum_{i=1}^n \nabla_{\theta_k} \ell(\theta_k, x_i) \kappa(\theta_k, \theta_j)}_{1. \text{ The Pull of the Data}} + \underbrace{\nabla_{\theta_k} \log \pi(\theta_k) \kappa(\theta_k, \theta_j)}_{2. \text{ The Anchor of the Prior}} + \underbrace{\nabla_{\theta_k} \kappa(\theta_k, \theta_j)}_{3. \text{ Kernel Repulsion}} \right]. \quad (1.5)$$

The first two terms represent the smoothed gradient (the score) of the generalised posterior $\nabla \log \pi_w$, acting as the driving force that drags particles towards high-density regions. In total, counting $\kappa(\theta_l, \theta)$ as the kernel-weighted average of the score (Also shown in Figure 1.1), this total attraction can be seen in one term $\kappa(\theta_l, \theta) \nabla_{\theta_l} \log p(\theta_l)$. The third term is the gradient of the kernel itself; which pushes neighboring particles apart and prevents mode collapse that ordinary gradient descent/ascent would produce. **The explicit deterministic repulsion is the defining mechanism of SVGD, contrasting it with other particle methods that rely on stochastic noise for dispersion.**

The cost of SVGD is the pairwise kernel evaluation, $O(M^2)$ per iteration, and, in common with kernel methods generally performance degrades in higher dimensions as the kernel's ability to discriminate between particles weakens.

'In this work we attempt to explain the variance collapse in SVGD. On the qualitative side, we compare the SVGD update with gradient descent on the maximum mean discrepancy (MMD) objective;' - very important for my work. From Understanding the Variance Collapse of SVGD in High Dimensions

1.3.2 Deep Ensembles: Omitting the Dispersion Force

Wild et al. (2023) (and Lakshminarayanan et al., 2017??? for specifically deep ensembles not as a bayesian thing) establish a rigorous link between Bayesian inference, variational methods, and deep ensembles. They achieve this by recasting the non-convex training objective of deep learning as a convex optimisation problem over the space of probability measures $\mathcal{P}(\Theta)$. This aligns perfectly with the optimisation-centric perspective introduced in Section 1, and they analyse this space using Wasserstein gradient flows. Concretely, they minimise the free energy:

$$L(q) = \int_{\Theta} V(\theta) dq(\theta) + \frac{\lambda_1}{2} \iint_{\Theta^2} \kappa(\theta, \theta') dq(\theta) dq(\theta') + \lambda_2 \int_{\Theta} q(\theta) \log q(\theta) d\theta, \quad \text{where } V(\theta) = \ell_n(\theta) - \lambda_2 \log \pi(\theta), \quad (1.6)$$

This objective can be broken down into three separate key ideas in a thermodynamic sum, similar to before. The first term is the external potential energy averaged over the particles where V is the potential. The potential energy $V = \ell_n - \lambda_2 \log \pi$ pulls the particles towards the data given by the loss function ℓ as well as the prior form $\lambda_2 \log \pi$. The second term, kernel interaction weighted by λ_1 , is similar to the SVGD interaction term. Finally, the third term is a negative-entropy term weighted by λ_2 .

Taking just the potential energy term and third term representing entropy, equivalent to setting λ_1 to 0, recovers our π_w from before. Starting with our total free energy equation,

$$\begin{aligned} L(q)|_{\lambda_1=0} &= \int_{\Theta} V(\theta) q(\theta) d\theta + \lambda_2 \int_{\Theta} q(\theta) \log q(\theta) d\theta \\ &= \int_{\Theta} (\ell_n(\theta) - \lambda_2 \log \pi(\theta)) q(\theta) d\theta + \lambda_2 \int_{\Theta} q(\theta) \log q(\theta) d\theta \\ &= \int_{\Theta} \ell_n(\theta) q(\theta) d\theta + \lambda_2 \int_{\Theta} q(\theta) (\log q(\theta) - \log \pi(\theta)) d\theta \\ &= \mathbb{E}_q[\ell_n(\theta)] + \lambda_2 \text{KL}(q||\pi). \end{aligned}$$

Dividing this objective by λ_2 exactly recovers the Rule of Three objective from Section 1, where the generalisation learning rate is $w = 1/\lambda_2$. Consequently, the minimiser of this free energy is the Gibbs posterior $q^* = \pi_w \propto \pi(\theta) \exp\{-\ell_n(\theta)/\lambda_2\}$. As established previously, in the specific case where $\lambda_2 = 1$ ($w = 1$) and the loss is the self-information $-\log p(x_i | \theta)$, this objective recovers the standard Bayesian posterior π_B .

Furthermore, we can see that when we set $\lambda_1 = \lambda_2 = 0$ we are left with only the $V(\theta)$ giving standard gradient descent and no dispersion of particles; the particles collapse onto the global minimiser of the loss. **this is a deep ensemble.**

Discretisation and Update Rule

As SVGD did and inline with particle approximations, to implement this infinite-dimensional gradient descent in practice, Wild et al. (2023) approximate the continuous-time Wasserstein gradient flow (**Show optimal transport and what the wasserstein gradient flow is**) using an ensemble of N_E interacting particles. The evolution of these particles in continuous time is governed by a stochastic differential equation (SDE) (**of McKean-Vlasov type**):

$$d\theta_n(t) = - \left(\nabla V(\theta_n(t)) + \frac{\lambda_1}{N_E} \sum_{j=1}^{N_E} \nabla_1 \kappa(\theta_n(t), \theta_j(t)) \right) dt + \sqrt{2\lambda_2} dB_n(t), \quad (1.7)$$

for $n = 1, \dots, N_E$, where $\{B_n(t)\}_{t>0}$ are independent standard Brownian motions, and $\nabla_1 \kappa$ denotes the gradient of the interaction kernel with respect to its first argument.

To simulate these trajectories computationally over a finite time horizon, the Euler-Maruyama discretisation (Appendix B in wild et al) method is applied. Given a discrete step size (or learn-

ing rate) $\eta > 0$, the update rule for the n -th particle at iteration t becomes

$$\theta_{n,t+1} = \theta_{n,t} - \eta \left(\nabla V(\theta_{n,t}) + \frac{\lambda_1}{N_E} \sum_{j=1}^{N_E} \nabla_1 \kappa(\theta_{n,t}, \theta_{j,t}) \right) + \sqrt{2\eta\lambda_2} Z_{n,t}. \quad (1.8)$$

where $Z_{n,t} \sim \mathcal{N}(0, I)$ ⁵ is standard Gaussian noise injected at each update step. Apparently. (still don't fully get this)

This update rule translates the thermodynamic properties of the free energy objective directly into algorithmic operations, recovering three distinct ensemble methodologies depending on the regularisation parameters:

- **Deep Ensembles (DE):** As established above, setting $\lambda_1 = \lambda_2 = 0$ removes both the repulsive interactions and the injected noise. The particles evolve independently via standard gradient descent from their random initialisations, gravitating towards the local minima of the loss landscape.
- **Deep Langevin Ensembles (DLE):** Setting $\lambda_1 = 0$ and $\lambda_2 > 0$ yields independent Langevin diffusions. The particles do not interact, but the injected noise ensures they sample from the target Gibbs posterior rather than simply collapsing onto a mode.
- **Deep Repulsive Langevin Ensembles (DRLE):** When both regularisation parameters are strictly positive ($\lambda_1 > 0, \lambda_2 > 0$), the full update rule is engaged. Particles actively repel each other via the kernel κ while exploring the landscape with injected noise, guaranteeing convergence to the global minimiser of the free energy over the space of probability measures.

Note here that while the particle motion that we showed above works well, similar to SGD, we can achieve speed increases by using methods such as ADAM or RMSprop with caution. Pre-conditioned SGLD (pSGLD) (Li et al., 2016) incorporates adaptive preconditioners (e.g., Adam or RMSprop) by scaling both the gradient trajectory and the injected noise inversely to the historical gradient variance, pSGLD allows particles to efficiently traverse ill-conditioned regions while maintaining the correct target posterior. Without this, the particles will not converge to the target posterior. (Welling and Teh, 2011)

Toy Example: The Global Minimiser

To explicitly demonstrate the theoretical convergence guarantees of these interacting particle systems, Wild et al. (2023) construct a one-dimensional loss landscape featuring a deep global minimum alongside a shallower local minimum. By initialising particles from a standard Gaussian prior, they show that unregularised Deep Ensembles merely partition particles into these

⁵The presence of η under the square root arises directly from the scaling properties of Brownian motion. When discretising the SDE, the infinitesimal increment $dB_n(t)$ is replaced by a finite time increment $\Delta B_n = B_n(t + \eta) - B_n(t)$. By definition, the increments of standard Brownian motion over a time step η are normally distributed with variance η , such that $\Delta B_n \sim \mathcal{N}(0, \eta I)$. To simulate this computationally using standard Gaussian noise $Z_{n,t} \sim \mathcal{N}(0, I)$, we must scale $Z_{n,t}$ by the standard deviation $\sqrt{\eta}$, yielding $\Delta B_n = \sqrt{\eta} Z_{n,t}$. Substituting this into the noise term of the SDE yields $\sqrt{2\lambda_2}(\sqrt{\eta} Z_{n,t}) = \sqrt{2\eta\lambda_2} Z_{n,t}$.

distinct modes based on their initial regions of attraction. In contrast, the noise and repulsive forces in fully regularised methods allow the ensemble to escape local traps and successfully represent the true global minimiser. The underlying non-convex loss function driving this behaviour is defined as an asymmetric polynomial:

$$\ell(\theta) = \frac{3}{2} \left(\frac{1}{4}\theta^4 + \frac{1}{3}\theta^3 - \theta^2 \right) - \frac{3}{8}, \quad \text{with prior } \pi(\theta) = \mathcal{N}(0, 1).$$

To visualise the distinct behaviour of each ensemble method within this asymmetric polynomial landscape, we simulate the particle trajectories over 100,000 steps. Figure 1.2 illustrates the progression from deterministic collapse to fully regularised exploration.




Figure 1.2: Particle distributions after 100,000 steps across three regularisation regimes. **(a)** Driven solely by gradient descent, unregularised particles collapse entirely into point masses at the local and global minima. **(b)** The injection of noise allows particles to escape strict collapse and sample from the optimal Gibbs posterior. **(c)** Adding a repulsive kernel forces particles away from one another, resulting in a broader, flatter distribution that ensures more diverse coverage of the parameter space.

While the Deep Repulsive Langevin Ensembles capture a much broader representation of the space, it is crucial to understand that increasing this repulsive force introduces a fundamental trade-off. Figure 1.3 provides a comprehensive analysis of how scaling λ_1 shifts the algorithmic behaviour away from the true, unregularised posterior.

1.3.3 MMD of *PrO*

Kernel mean embedding as

$$\mu_k(P) := \int k(\cdot, y) dP(y) \in \mathcal{H}(k)$$

Then the divergence of a candidate $P \in \mathcal{P}(\mathcal{Y})$ from the data-generating distribution $P_\star$ is given by the MMD

$$\text{MMD}(P_\star, P) = \|\mu_k(P_\star) - \mu_k(P)\|_{\mathcal{H}(k)}$$

- proper metrix if k is the characteristic kernel - Statistical divergence by "proper scoring rule [Dawid, 1986]"

Mixture model;

$$P_Q = \int P_\theta dQ(\theta) \in \mathcal{P}(\mathcal{Y})$$




Figure 1.3: Quantitative impact of the repulsive regularisation parameter λ_1 . As the repulsive force increases, probability mass systematically redistributes between basins, peak heights decrease, and the overall standard deviation widens. This increased particle diversity comes at the explicit cost of an expanding Kullback-Leibler (KL) divergence from the standard Deep Langevin Ensemble target.

the mean embedding for the P_Q is given as

$$\begin{aligned}
 \mu_k(P_Q) &= \int k(\cdot, y) dP_Q(y) \\
 &= \int k(\cdot, y) \int P_\theta dQ(\theta) \\
 &= \int \int k(\cdot, y) P_\theta dQ(\theta) \\
 &= \int \mu_k(P_\theta) dQ(\theta)
 \end{aligned}$$

which means that

$$MMD^2(P_\star, P) = \left\| \int \{\mu_k(P_\star) - \mu_k(P_\theta)\} dQ(\theta) \right\|_{\mathcal{H}(k)}^2 = \int \int \kappa_{P_\star}(\theta, \vartheta) dQ(\theta) dQ(\vartheta)$$

- last step is from expanding the norm using the inner product then exchanging inner product and integral

$$\begin{aligned}\left\|\int\{\dots\}dQ(\theta)\right\|^2 &= \langle \int(\dots)dQ(\theta), \int(\dots)dQ(\vartheta) \rangle \\ &= \int \int \langle \mu_k(P_\star) - \mu_k(P_\theta), \mu_k(P_\star) - \mu_k(P_\vartheta) \rangle_{\mathcal{H}(k)} dQ(\theta) dQ(\vartheta)\end{aligned}$$

where

$$\begin{aligned}\kappa_{P_\star}(\theta, \vartheta) &= \langle \mu_k(P_\star) - \mu_k(P_\theta), \mu_k(P_\star) - \mu_k(P_\vartheta) \rangle_{\mathcal{H}(k)} \\ &= \langle \mu_k(P_\star), \mu_k(P_\star) \rangle - \langle \mu_k(P_\star), \mu_k(P_\vartheta) \rangle - \langle \mu_k(P_\theta), \mu_k(P_\star) \rangle + \langle \mu_k(P_\theta), \mu_k(P_\vartheta) \rangle\end{aligned}$$

Gives us 4 terms

1. data vs data = constant
2. data vs model 1
3. data vs model 2
4. model 1 vs model 2

where

$$\kappa_{P_\star}(\theta, \vartheta) \stackrel{+C}{=} \iint k(y, y') dP_\theta(y) dP_\vartheta(y') - \iint k(y, y') dP_\theta(y) dP_\star(y') - \iint k(y, y') dP_\star(y) dP_\vartheta(y')$$

In practice we dont have the true data generating function $P_\star$ so we replace it with the empirical distribution $P_n = \frac{1}{n} \sum_i \delta_{y_i}$. Bellow the same as above but with the replaced P_n

$$\kappa_{P_n}(\theta, \vartheta) = \iint k(y, y') dP_\theta(y) dP_\vartheta(y') - \frac{1}{n} \sum_{i=1}^n \int k(y_i, y) dP_\theta(y) - \frac{1}{n} \sum_{i=1}^n \int k(y_i, y) dP_\vartheta(y)$$

$$\begin{aligned}\mathcal{F}_n(Q) &= \underbrace{\frac{1}{2} \iint \kappa_{P_n}(\theta, \vartheta) dQ(\theta) dQ(\vartheta)}_{\mathcal{E}_n(Q) \text{ (free energy)}} + \underbrace{\int v(\theta) dQ(\theta)}_{\text{single-particle (prior)}} + \underbrace{\lambda_n \int \log q(\theta) dQ(\theta)}_{\text{negative entropy}} \quad (1.9)\end{aligned}$$

$$= \frac{1}{2} \iint \kappa_{P_n}(\theta, \vartheta) dQ(\theta) dQ(\vartheta) + \lambda_n \text{KL}(Q \| Q_0) \quad (1.10)$$

Taking the gradint of everything for the particle update

$$\text{Functional Gradient:} \quad \nabla_W \mathcal{E}_n(Q_t)(\theta_t) = \nabla v(\theta_t) + \int \nabla_1 \kappa_{P_n}(\theta_t, \vartheta) dQ_t(\vartheta)$$

$$\text{Particle Update (SDE):} \quad d\theta_t = -\nabla_W \mathcal{E}_n(Q_t)(\theta_t) dt + \sqrt{2\lambda_n} dW_t$$

approximate Q_t with N particles $\theta_t^1, \dots, \theta_t^N$. The empirical distribution is:

$$Q_t \approx \hat{Q}_N^t = \frac{1}{N} \sum_{j=1}^N \delta_{\theta_t^j}$$

Now plug this into the integral:

$$\int \nabla_1 \kappa_{P_n}(\theta_t^i, \vartheta) dQ_t(\vartheta) \approx \int \nabla_1 \kappa_{P_n}(\theta_t^i, \vartheta) \frac{1}{N} \sum_{j=1}^N \delta_{\theta_t^j}(\vartheta)$$

The deltas collapse the integral into a sum. But you exclude $j = i$ (a particle shouldn't interact with itself), and compensate by using $N - 1$:

$$= \frac{1}{N-1} \sum_{j \neq i} \nabla_1 \kappa_{P_n}(\theta_t^i, \theta_t^j)$$

Substitute this back into the SDE, and give each particle its own independent Wiener process W_t^i :

$$d\theta_t^i = - \left(\nabla v(\theta_t^i) + \frac{1}{N-1} \sum_{j \neq i} \nabla_1 \kappa_{P_n}(\theta_t^i, \theta_t^j) \right) dt + \sqrt{2\lambda_n} dW_t^i$$

CHAPTER 2

Other notes for later chapters:

MoE Predictive

$$P_{MoE}(x) = \sum_{e=1}^E w_e(x) P_e(x) \quad (2.1)$$

$$= \sum_{e=1}^E w_e(x) \int P_\theta(x) dQ_e(\theta) \quad (2.2)$$

expanding out the MMD^2 term:

$$MMD^2(P, P_{MoE}) = \mathbb{E}_{x, x' \sim P} [k(x, x')] - 2\mathbb{E}_{x \sim P, x' \sim P_{MoE}} [k(x, x')] + \mathbb{E}_{x, x' \sim P_{MoE}} [k(x, x')]$$

$$T_{self} = \iint k(x, x') P_{MoE}(x) P_{MoE}(x') dx dx' \quad (2.3)$$

$$= \iint k(x, x') \left(\sum_{e=1}^E w_e(x) Q_e(x) \right) \left(\sum_{f=1}^E w_f(x') Q_f(x') \right) dx dx' \quad (2.4)$$

$$= \sum_{e=1}^E \sum_{f=1}^E \iint k(x, x') w_e(x) w_f(x') Q_e(x) Q_f(x') dx dx' \quad (2.5)$$

$$= \underbrace{\sum_{e=1}^E \iint k(x, x') w_e(x) w_e(x') Q_e(x) Q_e(x') dx dx'}_{\text{Intra-expert: } O(E \cdot n_e^2)} + \underbrace{\sum_{e \neq f} \iint k(x, x') w_e(x) w_f(x') Q_e(x) Q_f(x') dx dx'}_{\text{Inter-expert: } O(E^2 \cdot n_e^2)} \quad (2.6)$$

For the speed boost

$$\mathcal{C}_{ef} = \iint k(x, x') w_e(x) w_f(x') Q_e(x) Q_f(x') dx dx' = 0$$

KL divergence term

The global parameter distribution of the Mixture of Experts be the convex combination of the individual expert distributions:

$$Q_{MoE}(\theta) = \sum_{e=1}^E \pi_e Q_e(\theta)$$

(Where $\pi_e = \mathbb{E}_x[w_e(x)]$ is the expected routing weight for expert e across the dataset).

$$\begin{aligned} \mathcal{R}(Q_{1:E}) &= \text{KL}(Q_{MoE} || Q_0) = \int Q_{MoE}(\theta) \log \frac{Q_{MoE}(\theta)}{Q_0(\theta)} d\theta \\ &= \int \sum_{e=1}^E \pi_e Q_e(\theta) \log \left[\frac{\sum_{e=1}^E \pi_e Q_e(\theta)}{Q_0(\theta)} \right] d\theta \end{aligned}$$

Taking the derivative.

$$\frac{\delta \mathcal{R}(Q_{1:E})}{\delta Q_{MoE}(\theta)} = \log \frac{Q_{MoE}(\theta)}{Q_0(\theta)} + 1 \quad (2.7)$$

$$\frac{\delta Q_{MoE}(\theta)}{\delta Q_e(\theta)} = \frac{\delta}{\delta Q_e(\theta)} \sum_{e=1}^E \pi_e Q_e(\theta) = \pi_e \quad (2.8)$$

Combining (1) and (2) gives the general derivative before applying any MoE assumptions:

$$\begin{aligned} \frac{\delta \mathcal{R}}{\delta Q_e}(\theta) &= \pi_e \left(\log \frac{\sum_{f=1}^E \pi_f Q_f(\theta)}{Q_0(\theta)} + 1 \right) \\ &= \pi_e \left(\log \frac{\pi_e Q_e(\theta) + \sum_{f \neq e} \pi_f Q_f(\theta)}{Q_0(\theta)} + 1 \right) \\ &= \pi_e \left(\log \frac{Q_e(\theta)}{Q_0(\theta)} + \log \pi_e + 1 \right), \quad \text{if experts are non-overlapping.} \end{aligned}$$

Attempting to Isolate $Q_e(\theta)$ and solve for 0 to find the optimal..

$$\begin{aligned} 0 &= V_{Q_e}(\theta) + \lambda \pi_e \left(\log \frac{\pi_e Q_e(\theta) + \sum_{f \neq e} \pi_f Q_f(\theta)}{Q_0(\theta)} + 1 \right) + C \\ Q_e^*(\theta) &= \underbrace{\frac{1}{\pi_e Z_e} Q_0(\theta) \exp \left(-\frac{1}{\lambda \pi_e} V_{Q_e}(\theta) \right)}_{\text{Standard Gibbs Term}} - \underbrace{\sum_{f \neq e} \frac{\pi_f}{\pi_e} Q_f(\theta)}_{\text{Interference Term}} \end{aligned}$$

We now have the update gradients for each expert and globally:

$$\begin{aligned}\nabla_{\theta} \left(\frac{\delta \mathcal{R}}{\delta Q_e}(\theta) \right) &= \nabla_{\theta} \log \frac{Q_{MoE}(\theta)}{Q_0(\theta)} \\ &= \pi_e \nabla_{\theta} \log \left(\sum_{f=1}^E \pi_f Q_f(\theta) \right) - \pi_e \nabla_{\theta} \log Q_0(\theta)\end{aligned}\tag{2.9}$$

$$\nabla_{\theta} \left(\frac{\delta \mathcal{R}}{\delta Q_{MoE}}(\theta) \right) = \nabla_{\theta} \log Q_{MoE}(\theta) - \nabla_{\theta} \log Q_0(\theta)\tag{2.10}$$

Bibliography

- Aitchison, J. (1975). Goodness of prediction fit. *Biometrika*, 62(3):547–554.
- Bernardo, J. M. and Smith, A. F. (1994). *Bayesian Theory*. John Wiley & Sons, Chichester.
- Bissiri, P. G., Holmes, C. C., and Walker, S. G. (2016). A general framework for updating belief distributions. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 78(5):1103–1130.
- Blei, D. M., Kucukelbir, A., and McAuliffe, J. D. (2017). Variational inference: A review for statisticians. *Journal of the American Statistical Association*, 112(518):859–877.
- Gneiting, T. and Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378.
- Hoffman, M. D., Blei, D. M., Wang, C., and Paisley, J. (2013). Stochastic variational inference. *The Journal of Machine Learning Research*, 14(1):1303–1347.
- Knoblauch, J., Jewson, J., and Damoulas, T. (2022). An optimization-centric view on bayes’ rule: Reviewing and generalizing variational inference. *Journal of Machine Learning Research*, 23(132):1–109.
- Lai, J. and Yao, Y. (2024). Predictive variational inference: Learn the predictively optimal posterior distribution. *arXiv preprint arXiv:2410.14843*.
- Li, C., Chen, C., Carlson, D., and Carin, L. (2016). Preconditioned stochastic gradient langevin dynamics for deep neural networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 30.
- McLatchie, Y., Cherief-Abdellatif, B.-E., Frazier, D. T., and Knoblauch, J. (2025). Predictively oriented posteriors. *arXiv preprint arXiv:2510.01915*.
- Miller, J. W. (2021). Asymptotic normality, concentration, and coverage of generalized posteriors. *Journal of Machine Learning Research*, 22(168):1–53.
- Shen, Z., Knoblauch, J., Power, S., and Oates, C. J. (2024). Prediction-centric uncertainty quantification via mmd. *arXiv*.
- Welling, M. and Teh, Y. W. (2011). Bayesian learning via stochastic gradient langevin dynamics. In *Proceedings of the 28th International Conference on Machine Learning (ICML-11)*, pages 681–688.